

1 Bayesian Regression
 $w \sim N(0, \sigma_p^2 I)$, $\epsilon \sim N(0, \sigma_n^2 I)$, $y = Xw + \epsilon$
 $y|w \sim N(Xw, \sigma_n^2 I)$
 $w|y \sim N((X^T X + \lambda I)^{-1} X^T y, (X^T X + \lambda I)^{-1} \sigma_n^2)$

2 Kalman Filter
 $\begin{cases} X_{t+1} = F X_t + \epsilon_t & \epsilon_t \sim N(0, \Sigma_x) \\ Y_t = H X_t + \eta_t & \eta_t \sim N(0, \Sigma_y) \end{cases}$ $X_1 \sim N(\mu_p, \Sigma_p)$
Then if X_0 is Gaussian then $X_t|Y_{1:t} \sim N(\mu_t, \Sigma_t)$:
 $\mu_{t+1} = F \mu_t + K_{t+1}(y_{t+1} - H F \mu_t)$
 $\Sigma_{t+1} = (I - K_{t+1} H)(F \Sigma_t F^T + \Sigma_x)$
 $K_{t+1} = (F \Sigma_t F^T + \Sigma_x) H^T (H (F \Sigma_t F^T + \Sigma_x) H^T + \Sigma_y)^{-1}$

3 Gaussian Processes
 $f \sim GP(\mu, k) \Rightarrow \forall \{x_1, \dots, x_n\} \forall n < \infty$
 $[f(x_1) \dots f(x_n)] \sim N([\mu(x_1) \dots \mu(x_n)], K)$
where $K_{ij} = k(x_i, x_j)$

3.1 Gaussian Process Regression
 $f \sim GP(\mu, k)$ then: $f|y_{1:n}, x_{1:n} \sim GP(\tilde{\mu}, \tilde{k})$
 $\tilde{\mu}(x) = \mu(x) + K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} (y_A - \mu_A)$
 $\tilde{k}(x, x') = k(x, x') - K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} K_{A,x'}$
Where: $K_{A,x} = [k(x_1, x) \dots k(x_n, x)]^T$
 $[K_{AA}]_{ij} = k(x_i, x_j)$ and $\mu_A = [\mu(x_1) \dots \mu(x_n)]^T$

3.2 Kernels
 $k(x, y)$ is a kernel if it's symmetric semidefinite positive:
 $\forall \{x_1, \dots, x_n\}$ then for the Gram Matrix
 $[K]_{ij} = k(x_i, x_j)$ holds $c^T K c \geq 0 \forall c$
Some Kernels: (h is the bandwidth hyperp.)
Gaussian (rbf): $k(x, y) = \exp(-\frac{\|x-y\|^2}{h^2})$
Exponential: $k(x, y) = \exp(-\frac{\|x-y\|}{h})$
Linear kernel: $k(x, y) = x^T y$ (here $K_{AA} = X X^T$)

3.3 Optimization of Kernel Parameters
Given a dataset A , a kernel function $k(x, y; \theta)$.
 $y \sim N(0, K_y(\theta))$ where $K_y(\theta) = K_{AA}(\theta) + \sigma_n^2 I$
 $\hat{\theta} = \arg \max_{\theta} \log p(y|X; \theta)$
In GP: $\hat{\theta} = \arg \min_{\theta} y^T K_y^{-1}(\theta) y + \log |K_y(\theta)|$
We can from here $\nabla \downarrow$:
 $\nabla_{\theta} \log p(y|X; \theta) = \frac{1}{2} \text{tr}((\alpha \alpha^T - K^{-1}) \frac{\partial K}{\partial \theta})$, $\alpha = K^{-1} y$
Or we could also be bayesian about θ

3.4 Aproximation Techniques
Local method: $k(x_1, x_2) = 0$ if $\|x_1 - x_2\| \gg 1$

Random Fourier Features: if $k(x, y) = \kappa(x - y)$
 $p(w) = \mathcal{F}\{\kappa(\cdot), w\}$. Then $p(w)$ can be normalized to be a density.
 $\kappa(x - y) = \mathbb{E}_{p(w)} [\exp\{i w^T (x - y)\}]$ antitransform
 $\kappa(x - y) = \mathbb{E}_{b \sim \mathcal{U}([0, 2\pi]), w \sim p(w)} [z_{w,b}(x) z_{w,b}^*(y)]$

where $z_{w,b}(x) = \sqrt{2} \cos(w^T x + b)$. If can MC extract features z . If # features is $\ll n$ then this is faster ($X^T X$ vs $X X^T$)

Inducing points: We a vector of inducing variables u
 $f_A|u \sim N(K_{Au} K_{uu}^{-1} u, K_{AA} - K_{Au} K_{uu}^{-1} K_{uA})$
 $f_*|u \sim N(K_{*u} K_{uu}^{-1} u, K_{**} - K_{*u} K_{uu}^{-1} K_{u*})$

Subset of Regressors (SoR): $\blacksquare \rightarrow 0$
FITC: $\blacksquare \rightarrow$ its diagonal

4 Review of useful concepts and Introduction

4.1 Multivariate Gaussian
 $f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$
Suppose we have a Gaussian random vector
 $\begin{bmatrix} X_A \\ X_B \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu_A \\ \mu_B \end{bmatrix}, \begin{bmatrix} \Sigma_{AA} & \Sigma_{AB} \\ \Sigma_{BA} & \Sigma_{BB} \end{bmatrix}\right) \Rightarrow X_A|X_B = x_B \sim \mathcal{N}(\mu_A + \Sigma_{AB} \Sigma_{BB}^{-1} (x_B - \mu_B), \Sigma_{AA} - \Sigma_{AB} \Sigma_{BB}^{-1} \Sigma_{BA})$

4.2 Convex / Jensen's inequality
 $g(x)$ is convex $\Leftrightarrow x_1, x_2 \in \mathbb{R}, \lambda \in [0, 1] : g''(x) > 0$
 $g(\lambda x_1 + (1 - \lambda) x_2) \leq \lambda g(x_1) + (1 - \lambda) g(x_2)$
 $\varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)]$

4.3 Information Theory elements:
Entropy: $H(X) \doteq -\mathbb{E}_{x \sim p_X} [\log p_X(x)]$
 $H(X|Y) \doteq -\mathbb{E}_{(x,y) \sim p_{(X,Y)}} [\log p_{Y|X}(y|x)]$
if $X \sim \mathcal{N}(\mu, \Sigma) \Rightarrow H(X) = \frac{1}{2} \log[(2\pi e)^d \det(\Sigma)]$
Chain Rule: $H(X, Y) = H(Y|X) + H(X)$
Mutual Info: $I(X, Y) \doteq KL(p_{(X,Y)} \| p_X p_Y)$
 $I(X, Y) = H(X) - H(X|Y)$
if $X \sim \mathcal{N}(\mu, \Sigma)$, $Y = X + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$:
then $I(X, Y) = \frac{1}{2} \log[\det(I + \frac{1}{\sigma^2} \Sigma)]$

4.4 Kullback-Leiber divergence
 $KL(p||q) = \mathbb{E}_p[\log \frac{p(x)}{q(x)}]$
if $p_0 \sim \mathcal{N}(\mu_0, \Sigma_0)$, $p_1 \sim \mathcal{N}(\mu_1, \Sigma_1) \Rightarrow KL(p_0||p_1) = \frac{1}{2} (\text{tr}(\Sigma_1^{-1} \Sigma_0) + (\mu_1 - \mu_0)^T \Sigma_1^{-1} (\mu_1 - \mu_0) - k + \log \frac{|\Sigma_1|}{|\Sigma_0|})$
 $\hat{q} = \arg \min_q KL(p||q) \Rightarrow$ overconservative
 $\hat{q} = \arg \min_q KL(q||p) \Rightarrow$ overconfident

5 Approximate inference
5.1 Laplace Approximation
 $\hat{\theta} = \arg \max_{\theta} p(\theta|y)$
 $\Lambda = -\nabla_{\theta} \nabla_{\theta} \log p(\theta|y)|_{\theta=\hat{\theta}}$
 $p(\theta|y) \simeq q(\theta) = N(\hat{\theta}, \Lambda^{-1})$
5.2 Variational Inference
 $\hat{q} = \arg \min_{q \in \mathcal{Q}} KL(q||p(\cdot|y))$
 $\hat{q} = ELBO$ Evidence Lower Bound
 $ELBO \doteq \mathbb{E}_{\theta \sim q} [\log p(y|\theta)] - KL(q||p(\cdot)) \leq \log p(y)$
5.3 Markov Chain Monte Carlo
Idea: All we need is sampling from postirior

Ergodic Markov Chains:
 $\exists t \text{ s.t. } \mathbb{P}(i \rightarrow j \text{ in } t \text{ steps}) > 0 \quad \forall i, j \Rightarrow$
 $\exists ! \pi = \lim_{N \rightarrow \infty} \mathbb{P}(X_N = x)$ Limit distribution

Ergodic Theorem: if $(X_i)_{i \in \mathbb{N}}$ is ergodic:
 $\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N f(X_i) = \mathbb{E}_{x \sim \pi} [f(x)]$

Detailed Blanced Equation:
 $P(x|x')$ is the transition model of a MC:
if $R(x)P(x'|x) = R(x')P(x|x')$ then R is the limit distribution of the MC

Metropolis Hastings Algo: Sample from a MC which has $P(x) = \frac{Q(x)}{Z}$ as limit dist.

Result: $\{X_i\}_{i \in \mathbb{N}}$ sampled from the MC
init: $R(x|x')$
/* Good R choice \rightarrow fast convergence */
init: $X_0 = x_0$
for $t \leftarrow 1, 2, \dots$ **do**
 $x' \sim R(\cdot, x_{t-1})$
 $\alpha = \min\left\{1, \frac{Q(x')R(x_{t-1}|x')}{Q(x_{t-1})R(x'|x_{t-1})}\right\}$
 with probability α **do**
 $X_t = x'$;
 otherwise $X_t = x_{t-1}$;

Metropolis Adj. Langevin Algo (MALA):
Energy function: $P(x) = \frac{Q(x)}{Z} = \frac{1}{Z} \exp(-f(x))$
We chose: $R(x|x') = \mathcal{N}(x' - \tau \nabla f(x), 2\tau I)$
Stoch. Grad. Langevin Dynamics (SGLD):
We use SGD to Approximate ∇f . Converges also without acceptance step
Hamilton MC: SGD performance improved by adding momentum (consider last step ∇f)
Gibbs sampling: Practical when $X \in \mathbb{R}^n$
Used when $P(X_{1:n})$ is hard but $P(X_i|X_{-i})$ is easy.

init: $x_0 \in \mathbb{R}^n$; $(x_0^{(B)} = x^{(B)})$ B is our data
for $t = 1, 2, \dots$ **do**
 $x_t = x_{t-1}$
 with $i \sim \mathcal{U}(\{1 : n\} \setminus B)$ **do**
 $x_{t-1}^{(i)} \sim P(x^{(i)}|x^{(-i)})$

* if we do it $\forall i \notin B$ no DBE but more practical

5.4 Variable elimination for MPE (most probable explanation):

With loopy graphs, BP is often **overconfident/oscillates**.
6 Bayesian Neural Nets
Likelihood: $p(y|x; \theta) = \mathcal{N}(f_1(x, \theta), \exp(f_2(x, \theta)))$
Prior: $p(\theta) = \mathcal{N}(0, \sigma_p^2)$
 $\theta_{MAP} = \arg \max_{\theta} \log(p(y, \theta))$

6.1 Variational inference:
Usually we use $Q =$ Set of Gaussians
 $\hat{q} = \arg \max ELBO$ Reparameterization trick
 q approx. the posterior but how to predict?
 $p(y^*|x^*, \mathcal{D}) \simeq \frac{1}{m} \sum_{j=1}^m p(y^*|x^*, \theta^{(j)})$, $\theta \sim \hat{q}(\theta)$
Gaussian Mixture distribution: $\mathbb{V}(y^*|x^*, \mathcal{D}) \simeq \frac{1}{m} \sum_{j=1}^m \sigma^2(x^*, \theta^{(j)}) + \frac{1}{m} \sum_{j=1}^m (\mu(x^*, \theta^{(j)}) - \bar{\mu}(x^*))^2$
 $\blacksquare \rightarrow$ Alestoric, $\blacksquare \rightarrow$ Epistemic

Dropouts Regularization: Random ignore nodes in SGD iteration: Equavalent to VI with $Q = \{q(\cdot|\lambda) = \prod_j q_j(\theta_j|\lambda), \lambda \in \mathbb{R}^d\}$
where $q_j(\theta_j|\lambda) = p \delta_0(\theta_j) + (1 - p) \delta_{\lambda_j}(\theta_j)$
This allows to do Dropouts also in prediction

6.2 MCMC:
MCMC but cannot store all the $\theta^{(i)}$:
1) Subsampling: Only store a subset of the $\theta^{(i)}$:
2) Gaussian Aproximation: We only keep:
 $\mu_i = \frac{1}{T} \sum_{j=1}^T \theta_i^{(j)}$ and $\sigma_i = \frac{1}{T} \sum_{j=1}^T (\theta_i^{(j)} - \mu_i)^2$
And update them online.

Predictive Esnable NNs:
Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1:n}$ be our dataset.
Train θ_i^{MAP} on \mathcal{D}_i with $i = 1, \dots, m$
 \mathcal{D}_i is a Bootstrap of \mathcal{D} of same size
and $p(y^*|x^*, \mathcal{D}) \simeq \frac{1}{m} \sum_{j=1}^m p(y^*|x^*, \theta_i^{MAP})$

6.3 Model calibration
Train \hat{q} on \mathcal{D}_{train}
Evaluate \hat{q} on $\mathcal{D}_{val} = \{(y', x')\}_{i=1:m}$
Held-Out-Likelihood $\doteq \log p(y'_{1:m} | x'_{1:m}, \mathcal{D}_{train})$
 $\geq \mathbb{E}_{\theta \sim \hat{q}} [\sum_{i=1}^m \log p(y'_i | x'_i, \theta)]$ (Jensen)
 $\simeq \frac{1}{k} \sum_{j=1}^k \sum_{i=1}^m \log p(y'_i | x'_i, \theta^{(j)})$, $\theta^{(j)} \sim \hat{q}$

Evaluate predicted accuracy: We divide \mathcal{D}_{val} into bins according to predicted confidence values. In each bin we compare accuracy with confidence

7 Active Learning
Let \mathcal{D} be the set of observable points.
We can observe $\mathcal{S} \subseteq \mathcal{D}, |\mathcal{S}| \leq R$
Information Gain: $\hat{\mathcal{S}} = \arg \max_{\mathcal{S}} F(\mathcal{S}) = I(f, y_{\mathcal{S}})$
For GPs: $F(\mathcal{S}) = \frac{1}{2} \log |I + \frac{1}{\sigma^2} K_{\mathcal{S}\mathcal{S}}|$
This is NP Hard, \Rightarrow Greedy Algo:

init: $\mathcal{S}^* = \emptyset$
for $t = 1 : R$ **do**
 $x_t = \arg \max_{x \in \mathcal{D}} F(\mathcal{S}^* \cup \{x\})$
 $(x_t = \arg \max_{x \in \mathcal{D}} \sigma_x^2 | \mathcal{S} \text{ for GPs})$
 $(x_t = \arg \max_{x \in \mathcal{D}} \frac{\sigma_{f|S}^2(x)}{\sigma_n^2(x)} \text{ for heter. GPs})$
 $\mathcal{S}^* = \mathcal{S} \cup \{x_t\}$

F is **Submodular** if: $\forall x \in \mathcal{D}, \forall A \subseteq B \subseteq \mathcal{D}$ holds
that: $F(A \cup \{x\}) - F(A) \geq F(B \cup \{x\}) - F(B)$
F is Submodular $\Rightarrow F(\mathcal{S}^*) \geq \left(1 - \frac{1}{e}\right) F(\hat{\mathcal{S}})$

8 Bayesian Optimization

Like Active Learning but we only want to find the optima. We pick x_1, x_2, \dots from \mathcal{D} and observe $y_i = f(x_i) + \epsilon_i$.

Comulative regret: $R_T = \sum_{t=1}^T \left(\max_{x \in \mathcal{D}} f(x) - f(x_t) \right)$

Oss: $\frac{R_T}{T} \rightarrow 0 \Rightarrow \max_t f(x_t) \rightarrow \max_{x \in \mathcal{D}} f(x)$

8.1 Upper Confidence Sampling

With GP $x_t = \arg \max_{x \in \mathcal{D}} \mu_{t-1}(x) + \beta_t \sigma_{t-1}(x)$

Chosing the correct β_t we get: $\frac{R_T}{T} = \mathcal{O}\left(\sqrt{\frac{\gamma_T}{T}}\right)$.

Where $\gamma_t = \max_{|S| < T} I(f; y_S)$. On d dims:

Linear: $\gamma_T = \mathcal{O}(d \log T)$ RBF: $\gamma_T = \mathcal{O}((\log T)^{d+1})$

Optimal $\beta_t = \mathcal{O}(\|f\|_K^2 + \gamma_t \log^3 T)$

Oss: $\beta \uparrow$ = more exploration

8.2 Thompson Samling

$x_t = \arg \max_{x \in \mathcal{D}} \tilde{f}(x), \tilde{f} \sim p(f | x_{1:n}, y_{1:n})$

9 Markov Decision Process (MDP)